

Introduction to Parallel and Distributed Processing

Assignment 2

Author : Kavitha Elizebeth Varughese
UBID : 50560694

Overview

In this assignment, we design and implement an efficient MPI program for a fast application of sorting small integers. We will benchmark our solution to understand the scalability and performance of our parallel model over a sequential model.

Problem Description

Given an array of n small integers $X=[x_0, x_1, \dots, x_{n-1}]$ that is distributed among $p \ll n$ processors. Sort the array and distribute them among the processors such that each processor has none or a part of the array. We must ensure that $X_{li:l_j}$ and $X_{ki:k_j}$ distributed to processors with ranks l and k should be such that $X_{li:l_j} < X_{ki:k_j}$ if $l < k$.

For this problem, small integers are defined as 16-bit integers.

Proposed Solution

Our goal is to sort the small integers efficiently without a large all-to-all exchange of data. As we are dealing with small integers (there are only 65536 values), we will be performing a counting sort.

Streamlined Operations for Parallelization

1. In each processor, initialize a counting array of size 65536. These will be our buckets for our counting sort. In this array, we will count the number of occurrences of each value in X .

We must note that the values in X can be negative values as well. That's why, when we are incrementing the values in our counting array, we must add an OFFSET of 32768 to each value in X and increment for this index in our counting array. [APPENDIX 1.1]

2. Once we have the counting array of each processor ready, we can reduce the counting array across all the processors such that we have the cumulative count in processor 0. Now that we have our buckets full of all the values in X . We can broadcast the counting array to every processor for redistributing the data. We will be performing this operation using MPI constructs. [APPENDIX 1.2]
3. Finally, to redistribute the values, we will divide 65536 to p processors and ask each processor to expand their portion of the counting array and store the values in their local X .

Here, the values which are redistributed need not be uniformly distributed across the processors. It is important to resize the local array X to accommodate their portion of data. [APPENDIX 1.3]

Time Complexity

Sequential

We will only have 1 processor for a sequential model. The processor will have to go through all the values in X and then all the values in the counting array to reorder them . So the processor will visit every element exactly twice.

Therefore, the time complexity will be $O(n)$.

Parallel

All processors are only dealing with n/p values at a time. Reduce and broadcast will take $O(\log p)$. Thus, the overall time complexity should be $O(n/p) + \log(p)$

Performance Analysis

System Specification

We run the code across different nodes. The specifications of the nodes are as follows

Hosts used are in the cpn-f10-XX.core.ccr.buffalo.edu set.

Number of Cores	24
CPU	Intel(R) Xeon(R) Gold 6126 CPU @ 2.60GHz
Operating System	Ubuntu 24.04.1 LTS
Compiler	g++ (Gentoo 10.4.0 p5) 10.4.0
L1d cache	768 KiB

Input Specification

The code has been run on a single processor to record sequential time, followed by 8 nodes with 2/4/8/16 cores for 16/32/64/128 cores across nodes.

The code has been tested with a wide range of value of n to get a complete understanding of memory issues when it comes to parallel processing. The values are also selected to keep the ratio of n/p constant.

Results

N/P	1	16	32	64	128
10000000	0.0152988	0.0158068	0.00835738	0.00516979	0.00375179
160000000	0.225135	0.251104	0.132117	0.0658466	0.0360321
320000000	0.448639	0.499836	0.263452	0.129822	0.0658156
640000000	0.895514	0.998953	0.519678	0.259313	0.135098
1280000000	1.79503	1.99522	1.04248	0.509849	0.271325

N\P	1	16	32	64	128
5000	0.000561981	0.00121498	0.00108987	0.00149408	0.00168185
80000	0.000681464	0.00128306	0.00124299	0.00169662	0.00175185
160000	0.000780402	0.00194903	0.00129435	0.00162687	0.00172206
320000	0.00105284	0.00171584	0.00148921	0.00172987	0.00173508
640000	0.0014159	0.00204221	0.00166504	0.00231145	0.00178473

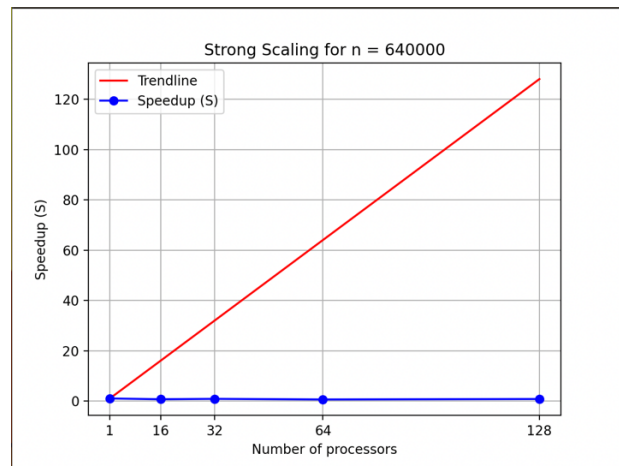
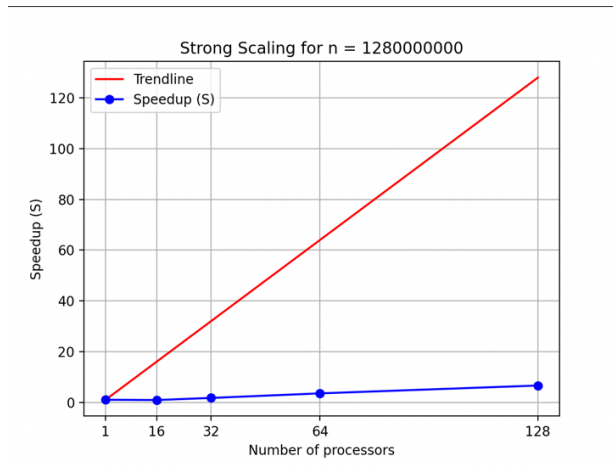
Strong Scaling Analysis

n = 12800000000

P	Speedup	Efficiency %
1	1.0	100.0
16	0.8996651998275880	5.622907498922420
32	1.7218843526974100	5.380888602179420
64	3.5207090726862300	5.501107926072230
128	6.615792868331340	5.168588178383860

n = 640000

P	Speedup	Efficiency %
1	1.0	100.0
16	0.6933175334564030	4.333234584102520
32	0.8503699610820160	2.6574061283813
64	0.6125592160764890	0.9571237751195140
128	0.7933412897188930	0.6197978825928850



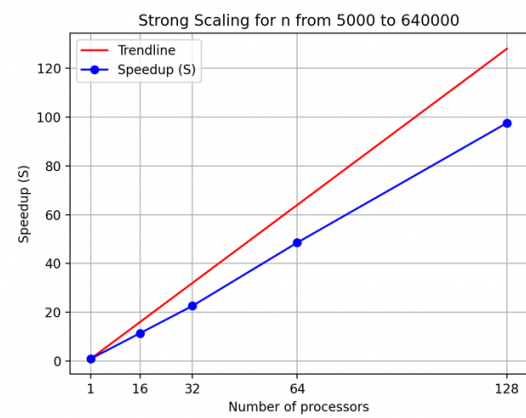
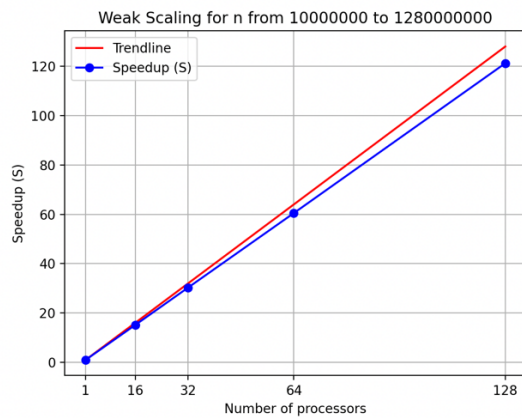
Weak Scaling Analysis

n from 10000000 to 12800000000

P	Speedup	Efficiency %
1	1.0	100.0
16	15.138590134938500	94.61618834336580
32	30.298613672140100	94.68316772543790
64	60.49022948030640	94.51598356297870
128	121.22126215617800	94.70411105951430

n from 5000 to 640000

P	Speedup	Efficiency %
1	1.0	100.0
16	11.431150310480900	71.44468944050570
32	22.615137602076400	70.67230500648860
64	48.55187400926150	75.86230313947110
128	97.58654602121860	76.23948907907710



Observation

The problem is not strongly scalable. This is because of the resizing of memory space in each process when redistributing the data. It is evident that there will exist at least one processor without any values in its local X. Thus, there will be processors who would need to allocate more memory to be able to accommodate the redistributed values. This extra need for memory highly affects the efficiency of our solution.

On the other hand, we can see that the problem is weakly scalable. We also see that the problem doesn't perform well for lower values of n. For higher values of n, the computation cost is probably high enough to overlook the memory costs. But, for lower values, the memory costs outshine computation costs, resulting in lesser efficiency.

With this problem, we can clearly understand the problems of memory allocations in parallel processing.

References

1. https://www.researchgate.net/publication/331221666_Parallel_Counting_Sort_A_Modified_of_Counting_Sort_Algorithm
2. <https://docs.ccr.buffalo.edu/en/latest/>

Appendix

- 1.1. Finding the count of every value of X locally in every processor.

```
// Each processor, count the occurrences of every number and save it in their local_count
std::vector<int> local_count(RANGE, 0);

for (const auto &x : Xi) {
    local_count[x + OFFSET]++;
}
```

- 1.2. Performing MPI Reduce and Broadcast to compute the global count of values in X across all processors.

```

std::vector<int> global_count(RANGE, 0);

// Combine all the local counts and find the global min and max
MPI_Reduce(local_count.data(), global_count.data(), RANGE, MPI_INT, MPI_SUM, 0, comm);

// Send global count to all processors as well as the global min and max
MPI_Bcast(global_count.data(), RANGE, MPI_INT, 0, comm);

```

1.3. Redistribution of data within every processor.

```

/*
   Divide global count vector equally to each processor between the min and max values
   Let each processor convert the counts back to the values and build Xi locally
*/

int start = rank * (RANGE / size);
int end = (rank + 1) * (RANGE / size);
int loc_n = std::accumulate(global_count.begin() + start, global_count.begin() + end, 0);

Xi.resize(loc_n);

// Convert global counts to Xi
int index = 0;
for (int i = start; i < end; i++) {
    while (global_count[i]-- > 0) {
        Xi[index++] = i - OFFSET;
    }
}

```